

Distributed Systems Module

Concurrency Parallelism & Distributed Systems

Carlos Segarra

January 8, 2020

Contents

1	Concepts of Distributed Systems	2
1.1	Definition of a Distributed System	2
1.2	Challenges of Distributed Systems	2
2	Distributed Algorithms	2
2.1	Time & Global States	2
2.2	Coordination and Agreement	3
3	Distributed Shared Data	3
3.1	Distributed Transactions	3
3.2	Consistency & Replication	3

1 Concepts of Distributed Systems

1.1 Definition of a Distributed System

1.2 Challenges of Distributed Systems

2 Distributed Algorithms

2.1 Time & Global States

Time

Physical Clocks

Physical clocks allow to synchronize nodes **within a given bound**. Synchronize at least every $R < \frac{\delta}{2\rho}$ **to limit skew between two clocks to less than δ** . Where:

- R : resynchronization interval.
- ρ : maximum clock drift rate.
- δ : maximum allowed clock skew.

Logical Clocks

Implemented to capture the happened-before relation. They satisfy:

1. If a and b are two events in the same process, $a \rightarrow b \Rightarrow C(a) < C(b)$
2. If a sends a message to $b \Rightarrow C(a) < C(b)$

Lamport's logical clocks:

- Each process P_i has a counter C_i
- C_i is updated using the following rules:
 1. When an event happens at P_i increment C_i by one.
 2. When P_i sends a message, set $ts(m) = C_i$
 3. When P_i receives a message, set $C_i = \max(C_i, ts(m))$ and then increase C_i by one.

Lamport's clocks do not guarantee that if $C(a) < C(b) \Rightarrow a \rightarrow b$.

Vector clocks:

- Each process P_i has an array $VC_i[1, \dots, n]$
- It is updated as follows:
 1. When P_i sends a message m , it adds 1 to $VC_i[i]$ and sends m with $ts(m) = VC_i$.
 2. When P_j receives a message from P_i , it updates each $VC_j[k]$ to $\max(VC_j[k], ts(m)[k])$ and increments $VC_j[j]$ by one.

Global States

A global state of the system is necessary for:

- Failure Recovery
- Detection of Properties: deadlocks, termination
- Debugging

We define some concepts:

1. The **history** of a process p is the sequence of events occurred at that process: $h(p) = \langle p_0, p_1, \dots \rangle$ (either internal or message sending).

2. The **state** i of process p is p 's history until event i : $s_i(p) = \langle p_0, \dots, p_i \rangle$.
3. The **global history** is the union of all the individual histories.
4. A **cut** is the global history up to a specific event in each process history.
5. A cut is **consistent** if it contains all the *happened-before* events. A consistent cut corresponds to a **consistent global state**.
6. A **run** is a total ordering of all events in a global history consistent with each local history.
7. A **linearization** or **consistent run** is a run consistent with the *happened-before* relation.
8. We say state S' is reachable from S if there is a linearization such that S precedes S' .

Distributed Snapshot

Global Predicates

Consistent global states form a lattice with reachability relation between sets. A **global state predicate**, φ is a property that is either true or false for a global state.

- A predicate is **stable** if once it becomes true, it remains true for all reachable states.
- A predicate is **non-stable** if it can become true and then false.
- A predicate φ possibly happened: if it is true for any of the consistent states in the lattice.
- A predicate φ definitely happened: if all paths from origin to end contain a consistent global state for which the predicate is true.

2.2 Coordination and Agreement

Mutual Exclusion

Election Algorithms

Multicast Communications

Consensus

3 Distributed Shared Data

3.1 Distributed Transactions

Introduction

Problems with Concurrent Transactions

Concurrency Control

Distributed Transactions

3.2 Consistency & Replication

Introduction

Data-Centric Consistency Models

Client-Centric Consistency Models

Consistency Protocols